

mlp-for-classification

May 4, 2024

1 MLP for Classification

1.1 Imports

```
[1]: import pandas as pd
import matplotlib.pyplot as plt
import torch
from torch import nn, optim
import torch.nn.functional as F
from torch.utils.data import DataLoader, TensorDataset
from sklearn.metrics import r2_score
from sklearn.model_selection import train_test_split
```

1.2 Load Data

```
[2]: !cp "/content/drive/MyDrive/colab_projects/MLP for Classification/mobile_price.
↪zip" /content
!unzip "/content/drive/MyDrive/colab_projects/MLP for Classification/
↪mobile_price.zip"
```

```
Archive: /content/drive/MyDrive/colab_projects/MLP for
Classification/mobile_price.zip
  inflating: test.csv
  inflating: train.csv
```

```
[3]: df = pd.read_csv('/content/train.csv')
df.head()
```

```
[3]:   battery_power  blue  clock_speed  dual_sim  fc  four_g  int_memory  m_dep \
0           842     0         2.2         0    1      0           7    0.6
1          1021     1         0.5         1    0      1          53    0.7
2           563     1         0.5         1    2      1          41    0.9
3           615     1         2.5         0    0      0          10    0.8
4          1821     1         1.2         0   13      1          44    0.6

   mobile_wt  n_cores  ...  px_height  px_width  ram  sc_h  sc_w  talk_time \
0         188        2  ...         20       756  2549    9     7          19
1         136        3  ...         905      1988  2631   17     3           7
```

2	145	5	...	1263	1716	2603	11	2	9
3	131	6	...	1216	1786	2769	16	8	11
4	141	2	...	1208	1212	1411	8	2	15

	three_g	touch_screen	wifi	price_range
0	0	0	1	1
1	1	1	0	2
2	1	1	0	2
3	1	0	0	2
4	1	1	0	1

[5 rows x 21 columns]

```
[4]: X = df.drop('price_range', axis = 1)
```

```
[5]: y = df['price_range']
```

```
[6]: x_train, x_valid, y_train, y_valid = train_test_split(X, y, test_size=0.7,
↳ random_state=42)
```

1.3 Preprocessing

1.3.1 Getting info about df

```
[7]: df.shape
```

```
[7]: (2000, 21)
```

```
[8]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2000 entries, 0 to 1999
Data columns (total 21 columns):
#   Column          Non-Null Count  Dtype
---  -
0   battery_power    2000 non-null   int64
1   blue             2000 non-null   int64
2   clock_speed      2000 non-null   float64
3   dual_sim         2000 non-null   int64
4   fc               2000 non-null   int64
5   four_g           2000 non-null   int64
6   int_memory       2000 non-null   int64
7   m_dep            2000 non-null   float64
8   mobile_wt        2000 non-null   int64
9   n_cores          2000 non-null   int64
10  pc               2000 non-null   int64
11  px_height        2000 non-null   int64
```

```

12  px_width      2000 non-null  int64
13  ram           2000 non-null  int64
14  sc_h          2000 non-null  int64
15  sc_w          2000 non-null  int64
16  talk_time     2000 non-null  int64
17  three_g       2000 non-null  int64
18  touch_screen  2000 non-null  int64
19  wifi          2000 non-null  int64
20  price_range   2000 non-null  int64
dtypes: float64(2), int64(19)
memory usage: 328.2 KB

```

1.3.2 Convert to Tensor

Note that since it's a classification problem, torch expects an integer as output and not a float.

```
[9]: x_train = torch.FloatTensor(x_train.values)
     y_train = torch.LongTensor(y_train.values)
```

```
[10]: x_valid = torch.FloatTensor(x_valid.values)
      y_valid = torch.LongTensor(y_valid.values)
```

1.3.3 Standardization

```
[11]: mu = x_train.mean(dim=0)
      std = x_train.std(dim=0)
```

```
[12]: x_train = (x_train - mu) / std
      x_valid = (x_valid - mu) / std
```

1.4 Dataloader

```
[13]: train_data = TensorDataset(x_train, y_train)
      train_loader = DataLoader(train_data, batch_size=100, shuffle=True)
```

```
[14]: valid_data = TensorDataset(x_valid, y_valid)
      valid_loader = DataLoader(valid_data, batch_size=200, shuffle=True)
```

1.5 Model

```
[15]: num_feats = 20
      num_class = 4
      h1 = 64
      h2 = 32

      model = nn.Sequential(nn.Linear(num_feats, h1),
                           nn.ReLU(),
```

```
nn.Linear(h1, h2),
nn.ReLU(),
nn.Linear(h2, num_class))
```

1.6 Device

```
[16]: device = 'cuda' if torch.cuda.is_available() else 'cpu'
```

```
[17]: model = model.to(device)
```

1.7 Loss and Optimizer

```
[18]: loss_fn = nn.CrossEntropyLoss()
optimizer = optim.SGD(model.parameters(), lr=0.01)
```

```
[19]: class AverageMeter(object):
    """Computes and stores the average and current value"""
    def __init__(self):
        self.reset()

    def reset(self):
        self.val = 0
        self.avg = 0
        self.sum = 0
        self.count = 0

    def update(self, val, n=1):
        self.val = val
        self.sum += val * n
        self.count += n
        self.avg = self.sum / self.count
```

1.8 Torchmetrics

In order to see how well our model has predicted the corresponding classes of each sample, we can use the accuracy metric in torchmetrics

```
[20]: !pip install torchmetrics
from torchmetrics import Accuracy
```

Collecting torchmetrics

Downloading torchmetrics-1.3.2-py3-none-any.whl (841 kB)

841.5/841.5

kB 6.0 MB/s eta 0:00:00

Requirement already satisfied: numpy>1.20.0 in

/usr/local/lib/python3.10/dist-packages (from torchmetrics) (1.25.2)

Requirement already satisfied: packaging>17.1 in /usr/local/lib/python3.10/dist-packages (from torchmetrics) (24.0)

Requirement already satisfied: torch>=1.10.0 in /usr/local/lib/python3.10/dist-packages (from torchmetrics) (2.2.1+cu121)

Collecting lightning-utilities>=0.8.0 (from torchmetrics)

 Downloading lightning_utilities-0.11.2-py3-none-any.whl (26 kB)

Requirement already satisfied: setuptools in /usr/local/lib/python3.10/dist-packages (from lightning-utilities>=0.8.0->torchmetrics) (67.7.2)

Requirement already satisfied: typing-extensions in /usr/local/lib/python3.10/dist-packages (from lightning-utilities>=0.8.0->torchmetrics) (4.11.0)

Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-packages (from torch>=1.10.0->torchmetrics) (3.14.0)

Requirement already satisfied: sympy in /usr/local/lib/python3.10/dist-packages (from torch>=1.10.0->torchmetrics) (1.12)

Requirement already satisfied: networkx in /usr/local/lib/python3.10/dist-packages (from torch>=1.10.0->torchmetrics) (3.3)

Requirement already satisfied: jinja2 in /usr/local/lib/python3.10/dist-packages (from torch>=1.10.0->torchmetrics) (3.1.3)

Requirement already satisfied: fsspec in /usr/local/lib/python3.10/dist-packages (from torch>=1.10.0->torchmetrics) (2023.6.0)

Collecting nvidia-cuda-nvrtc-cu12==12.1.105 (from torch>=1.10.0->torchmetrics)

 Using cached nvidia_cuda_nvrtc_cu12-12.1.105-py3-none-manylinux1_x86_64.whl (23.7 MB)

Collecting nvidia-cuda-runtime-cu12==12.1.105 (from torch>=1.10.0->torchmetrics)

 Using cached nvidia_cuda_runtime_cu12-12.1.105-py3-none-manylinux1_x86_64.whl (823 kB)

Collecting nvidia-cuda-cupti-cu12==12.1.105 (from torch>=1.10.0->torchmetrics)

 Using cached nvidia_cuda_cupti_cu12-12.1.105-py3-none-manylinux1_x86_64.whl (14.1 MB)

Collecting nvidia-cudnn-cu12==8.9.2.26 (from torch>=1.10.0->torchmetrics)

 Using cached nvidia_cudnn_cu12-8.9.2.26-py3-none-manylinux1_x86_64.whl (731.7 MB)

Collecting nvidia-cublas-cu12==12.1.3.1 (from torch>=1.10.0->torchmetrics)

 Using cached nvidia_cublas_cu12-12.1.3.1-py3-none-manylinux1_x86_64.whl (410.6 MB)

Collecting nvidia-cufft-cu12==11.0.2.54 (from torch>=1.10.0->torchmetrics)

 Using cached nvidia_cufft_cu12-11.0.2.54-py3-none-manylinux1_x86_64.whl (121.6 MB)

Collecting nvidia-curand-cu12==10.3.2.106 (from torch>=1.10.0->torchmetrics)

 Using cached nvidia_curand_cu12-10.3.2.106-py3-none-manylinux1_x86_64.whl (56.5 MB)

Collecting nvidia-cusolver-cu12==11.4.5.107 (from torch>=1.10.0->torchmetrics)

 Using cached nvidia_cusolver_cu12-11.4.5.107-py3-none-manylinux1_x86_64.whl (124.2 MB)

Collecting nvidia-cuspars-cu12==12.1.0.106 (from torch>=1.10.0->torchmetrics)

 Using cached nvidia_cuspars-cu12-12.1.0.106-py3-none-manylinux1_x86_64.whl (196.0 MB)

Collecting nvidia-nccl-cu12==2.19.3 (from torch>=1.10.0->torchmetrics)
 Using cached nvidia_nccl_cu12-2.19.3-py3-none-manylinux1_x86_64.whl (166.0 MB)
 Collecting nvidia-nvtx-cu12==12.1.105 (from torch>=1.10.0->torchmetrics)
 Using cached nvidia_nvtx_cu12-12.1.105-py3-none-manylinux1_x86_64.whl (99 kB)
 Requirement already satisfied: triton==2.2.0 in /usr/local/lib/python3.10/dist-packages (from torch>=1.10.0->torchmetrics) (2.2.0)
 Collecting nvidia-nvjitlink-cu12 (from nvidia-cusolver-cu12==11.4.5.107->torch>=1.10.0->torchmetrics)
 Using cached nvidia_nvjitlink_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl (21.1 MB)
 Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.10/dist-packages (from jinja2->torch>=1.10.0->torchmetrics) (2.1.5)
 Requirement already satisfied: mpmath>=0.19 in /usr/local/lib/python3.10/dist-packages (from sympy->torch>=1.10.0->torchmetrics) (1.3.0)
 Installing collected packages: nvidia-nvtx-cu12, nvidia-nvjitlink-cu12, nvidia-nccl-cu12, nvidia-curand-cu12, nvidia-cufft-cu12, nvidia-cuda-runtime-cu12, nvidia-cuda-nvrtc-cu12, nvidia-cuda-cupti-cu12, nvidia-cublas-cu12, lightning-utilities, nvidia-cuspars-cu12, nvidia-cudnn-cu12, nvidia-cusolver-cu12, torchmetrics
 Successfully installed lightning-utilities-0.11.2 nvidia-cublas-cu12-12.1.3.1 nvidia-cuda-cupti-cu12-12.1.105 nvidia-cuda-nvrtc-cu12-12.1.105 nvidia-cuda-runtime-cu12-12.1.105 nvidia-cudnn-cu12-8.9.2.26 nvidia-cufft-cu12-11.0.2.54 nvidia-curand-cu12-10.3.2.106 nvidia-cusolver-cu12-11.4.5.107 nvidia-cuspars-cu12-12.1.0.106 nvidia-nccl-cu12-2.19.3 nvidia-nvjitlink-cu12-12.4.127 nvidia-nvtx-cu12-12.1.105 torchmetrics-1.3.2

1.9 Model Training

```
[21]: num_epochs = 400

loss_train_hist = []
loss_valid_hist = []

acc_train_hist = []
acc_valid_hist = []

for epoch in range(num_epochs):
    loss_train = AverageMeter()
    acc_train = Accuracy(task='multiclass', num_classes=4).to(device)
    for i, (inputs, targets) in enumerate(train_loader):
        inputs = inputs.to(device)
        targets = targets.to(device)

        outputs = model(inputs)

        loss = loss_fn(outputs, targets)
```

```

    loss.backward()
    optimizer.step()
    optimizer.zero_grad()

    loss_train.update(loss.item())
    acc_train(outputs, targets)

    with torch.no_grad():
        loss_valid = AverageMeter()
        acc_valid = Accuracy(task='multiclass', num_classes=4).to(device)
        for i, (inputs, targets) in enumerate(valid_loader):
            inputs = inputs.to(device)
            targets = targets.to(device)

            outputs = model(inputs)
            loss = loss_fn(outputs, targets)

            loss_valid.update(loss.item())
            acc_valid(outputs, targets)

    loss_train_hist.append(loss_train.avg)
    loss_valid_hist.append(loss_valid.avg)

    acc_train_hist.append(acc_train.compute())
    acc_valid_hist.append(acc_valid.compute())

    if epoch % 10 == 0:
        print(f'Epoch {epoch}')
        print(f'Train: Loss = {loss_train.avg:.4}, Acc = {acc_train.compute():.4}')
        print(f'Valid: Loss = {loss_valid.avg:.4}, Acc = {acc_valid.compute():.4}')
        print()

```

Epoch 0

Train: Loss = 1.394, Acc = 0.2367

Valid: Loss = 1.394, Acc = 0.2486

Epoch 10

Train: Loss = 1.381, Acc = 0.2517

Valid: Loss = 1.382, Acc = 0.2571

Epoch 20

Train: Loss = 1.369, Acc = 0.2783

Valid: Loss = 1.371, Acc = 0.2779

Epoch 30

Train: Loss = 1.356, Acc = 0.3167

Valid: Loss = 1.358, Acc = 0.3121

Epoch 40

Train: Loss = 1.34, Acc = 0.39

Valid: Loss = 1.342, Acc = 0.3714

Epoch 50

Train: Loss = 1.32, Acc = 0.46

Valid: Loss = 1.322, Acc = 0.4421

Epoch 60

Train: Loss = 1.293, Acc = 0.4883

Valid: Loss = 1.295, Acc = 0.4793

Epoch 70

Train: Loss = 1.258, Acc = 0.5317

Valid: Loss = 1.26, Acc = 0.5121

Epoch 80

Train: Loss = 1.213, Acc = 0.5467

Valid: Loss = 1.215, Acc = 0.5229

Epoch 90

Train: Loss = 1.157, Acc = 0.5583

Valid: Loss = 1.16, Acc = 0.5329

Epoch 100

Train: Loss = 1.093, Acc = 0.5633

Valid: Loss = 1.096, Acc = 0.5486

Epoch 110

Train: Loss = 1.026, Acc = 0.5883

Valid: Loss = 1.029, Acc = 0.5643

Epoch 120

Train: Loss = 0.9583, Acc = 0.61

Valid: Loss = 0.9624, Acc = 0.5829

Epoch 130

Train: Loss = 0.8943, Acc = 0.6433

Valid: Loss = 0.899, Acc = 0.6129

Epoch 140

Train: Loss = 0.8351, Acc = 0.6733

Valid: Loss = 0.8405, Acc = 0.6586

Epoch 150

Train: Loss = 0.7809, Acc = 0.725

Valid: Loss = 0.7872, Acc = 0.7

Epoch 160

Train: Loss = 0.7314, Acc = 0.775

Valid: Loss = 0.7387, Acc = 0.7429

Epoch 170

Train: Loss = 0.6859, Acc = 0.8083

Valid: Loss = 0.6946, Acc = 0.7743

Epoch 180

Train: Loss = 0.6443, Acc = 0.8333

Valid: Loss = 0.6546, Acc = 0.8029

Epoch 190

Train: Loss = 0.6062, Acc = 0.8483

Valid: Loss = 0.6183, Acc = 0.8214

Epoch 200

Train: Loss = 0.5711, Acc = 0.87

Valid: Loss = 0.5852, Acc = 0.84

Epoch 210

Train: Loss = 0.5389, Acc = 0.8883

Valid: Loss = 0.555, Acc = 0.8593

Epoch 220

Train: Loss = 0.509, Acc = 0.9033

Valid: Loss = 0.5274, Acc = 0.8686

Epoch 230

Train: Loss = 0.4814, Acc = 0.91

Valid: Loss = 0.502, Acc = 0.8736

Epoch 240

Train: Loss = 0.4558, Acc = 0.9167

Valid: Loss = 0.4787, Acc = 0.8843

Epoch 250

Train: Loss = 0.4319, Acc = 0.9167

Valid: Loss = 0.4574, Acc = 0.8914

Epoch 260

Train: Loss = 0.4097, Acc = 0.92

Valid: Loss = 0.4378, Acc = 0.8986

Epoch 270

Train: Loss = 0.3889, Acc = 0.9233

Valid: Loss = 0.4197, Acc = 0.9057

Epoch 280

Train: Loss = 0.37, Acc = 0.9317

Valid: Loss = 0.403, Acc = 0.9086

Epoch 290

Train: Loss = 0.3519, Acc = 0.9383

Valid: Loss = 0.3877, Acc = 0.9086

Epoch 300

Train: Loss = 0.335, Acc = 0.9433

Valid: Loss = 0.3734, Acc = 0.9107

Epoch 310

Train: Loss = 0.3193, Acc = 0.9467

Valid: Loss = 0.3604, Acc = 0.9107

Epoch 320

Train: Loss = 0.3045, Acc = 0.9517

Valid: Loss = 0.3482, Acc = 0.9114

Epoch 330

Train: Loss = 0.2907, Acc = 0.9533

Valid: Loss = 0.337, Acc = 0.9129

Epoch 340

Train: Loss = 0.2778, Acc = 0.9567

Valid: Loss = 0.3267, Acc = 0.9114

Epoch 350

Train: Loss = 0.2657, Acc = 0.9617

Valid: Loss = 0.3172, Acc = 0.9114

Epoch 360

Train: Loss = 0.2542, Acc = 0.9633

Valid: Loss = 0.3084, Acc = 0.9136

Epoch 370

Train: Loss = 0.2435, Acc = 0.965

Valid: Loss = 0.3002, Acc = 0.9136

Epoch 380

Train: Loss = 0.2334, Acc = 0.9667

Valid: Loss = 0.2926, Acc = 0.9157

Epoch 390

Train: Loss = 0.2238, Acc = 0.9683

Valid: Loss = 0.2856, Acc = 0.9157


1.10 Evaluation

1.10.1 Loss

```
[22]: plt.plot(range(num_epochs), loss_train_hist, 'r-', label='Train')
plt.plot(range(num_epochs), loss_valid_hist, 'b-', label='Validation')

plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.grid(True)
plt.legend()
```

[22]: <matplotlib.legend.Legend at 0x7ed8b6051900>

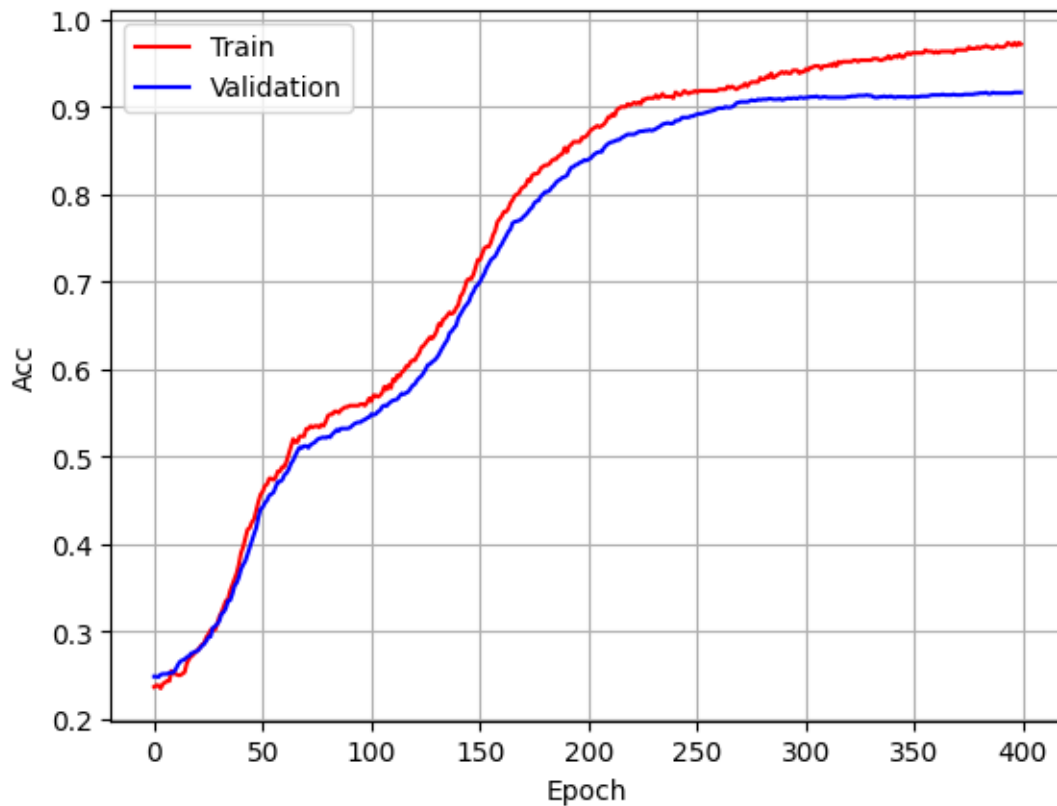


1.10.2 Accuracy

```
[23]: plt.plot(range(num_epochs), acc_train_hist, 'r-', label='Train')
plt.plot(range(num_epochs), acc_valid_hist, 'b-', label='Validation')

plt.xlabel('Epoch')
plt.ylabel('Acc')
plt.grid(True)
plt.legend()
```

[23]: <matplotlib.legend.Legend at 0x7ed8b6052e60>



1.11 Save Model

```
[24]: torch.save(model, 'model.pth')
```

```
[25]: mymodel = torch.load('model.pth')
```