

mlp-for-regression

May 4, 2024

1 MLP for Regression

1.1 Imports

```
[1]: import pandas as pd
import matplotlib.pyplot as plt
import torch
from torch import nn, optim
import torch.nn.functional as F
from torch.utils.data import DataLoader, TensorDataset
from sklearn.metrics import r2_score
```

1.2 Load Data

```
[2]: !cp "/content/drive/MyDrive/colab_projects/MLP for Regression/test.csv" test.csv
!cp "/content/drive/MyDrive/colab_projects/MLP for Regression/train.csv" train.
↪ csv
```

```
[3]: df_train = pd.read_csv('/content/train.csv')
df_train.head()
```

```
[3]:
```

	MedInc	HouseAge	AveRooms	AveBedrms	Population	AveOccup	Latitude	\
0	1.4817	6	4.443645	1.134293	1397	3.350120	36.77	
1	6.9133	8	5.976471	1.026471	862	2.535294	33.68	
2	1.5536	25	4.088785	1.000000	931	4.350467	36.60	
3	1.5284	31	2.740088	1.008811	597	2.629956	34.10	
4	4.0815	21	5.166667	1.002688	1130	3.037634	37.79	

	Longitude	MedHouseVal
0	-119.84	0.720
1	-117.80	2.741
2	-120.19	0.583
3	-118.32	2.000
4	-121.23	1.179

```
[4]: x_train = df_train.drop('MedHouseVal', axis=1)
y_train = df_train['MedHouseVal']
```

```
[5]: df_test = pd.read_csv('/content/test.csv')
df_test.head()
```

```
[5]:   MedInc  HouseAge  AveRooms  AveBedrms  Population  AveOccup  Latitude  \
0  6.4114      15  7.527559   1.049869      2529  3.318898    33.60
1  4.1843      12  6.330084   1.041783      2033  2.831476    38.62
2  1.7411      35  5.369159   1.294393       909  4.247664    33.93
3  3.4412      39  4.173405   1.074573      2156  1.937107    37.59
4  7.8195      16  8.602349   1.058725      2142  3.593960    33.70

      Longitude  MedHouseVal
0    -117.65      2.787
1    -120.91      2.076
2    -118.23      0.967
3    -122.37      3.538
4    -117.98      3.905
```

```
[6]: x_test = df_test.drop('MedHouseVal', axis=1)
y_test = df_test['MedHouseVal']
```

1.3 Preprocessing

1.3.1 Getting info about df

```
[7]: df_train.shape, df_test.shape
```

```
[7]: ((16512, 9), (4128, 9))
```

```
[8]: df_train.info
```

```
[8]: <bound method DataFrame.info of
Population  AveOccup  Latitude  \
0      1.4817      6  4.443645   1.134293      1397  3.350120    36.77
1      6.9133      8  5.976471   1.026471      862  2.535294    33.68
2      1.5536     25  4.088785   1.000000      931  4.350467    36.60
3      1.5284     31  2.740088   1.008811      597  2.629956    34.10
4      4.0815     21  5.166667   1.002688     1130  3.037634    37.79
...      ...      ...      ...      ...      ...      ...
16507  3.0625      20  5.860000   1.112000      745  2.980000    39.59
16508  2.6133     49  5.163755   1.100437     1131  2.469432    38.11
16509  4.4958     19  5.899767   1.074592     1206  2.811189    38.92
16510  2.5750     39  3.591203   1.086675     2546  3.293661    34.16
16511  2.2478     31  5.123810   1.100000     1259  2.997619    40.80

      Longitude  MedHouseVal
0    -119.84      0.720
1    -117.80      2.741
```

2	-120.19	0.583
3	-118.32	2.000
4	-121.23	1.179
...
16507	-121.90	0.938
16508	-122.25	1.031
16509	-121.22	1.926
16510	-118.14	1.535
16511	-124.13	0.811

[16512 rows x 9 columns]>

1.4 Convert to Tensor

```
[9]: x_train = torch.FloatTensor(x_train.values)
     y_train = torch.FloatTensor(y_train.values)
```

```
[10]: x_test = torch.FloatTensor(x_test.values)
      y_test = torch.FloatTensor(y_test.values)
```

1.5 Standardization

```
[11]: mu = x_train.mean(axis=0)
      std = x_train.std(axis=0)
      x_train = (x_train - mu) / std
```

```
[12]: x_test = (x_test - mu) / std
```

1.6 Dataloader

```
[13]: train_dataset = TensorDataset(x_train, y_train)
      train_loader = DataLoader(train_dataset, batch_size = 128, shuffle = True)
```

```
[14]: for x_batch, y_batch in train_loader:
      print(x_batch.shape, y_batch.shape)
      break
```

torch.Size([128, 8]) torch.Size([128])

```
[15]: test_dataset = TensorDataset(x_test, y_test)
      test_loader = DataLoader(test_dataset, batch_size = 258, shuffle = False)
```

1.7 Model

```
[16]: num_feats = 8
      out_feat = 1
      h1 = 64
      h2 = 32

      model = nn.Sequential(nn.Linear(num_feats, h1),
                            nn.ReLU(),
                            nn.Linear(h1, h2),
                            nn.ReLU(),
                            nn.Linear(h2, out_feat))
```

1.8 Loss and Optimizer

```
[17]: loss_fn = nn.MSELoss()
      optimizer = optim.SGD(model.parameters(), lr=0.001)
```

1.9 Model Training

```
[18]: num_epochs = 400

      for epoch in range(num_epochs):
          loss_train = 0
          for x_batch, y_batch in train_loader:
              yp = model(x_batch)
              loss = loss_fn(yp.squeeze(), y_batch)
              loss_train += loss
              loss.backward()
              optimizer.step()
              optimizer.zero_grad()

          loss_test = 0
          for x_batch, y_batch in test_loader:
              yp = model(x_batch)
              loss_test += loss_fn(yp.squeeze(), y_batch)

          if epoch % 10 == 0:
              loss_total_train = loss_train/len(train_loader)
              loss_total_test = loss_test/len(test_loader)
              print(f'Epoch = {epoch} : Loss Train = {loss_total_train:.4}')
              print(f'Epoch = {epoch} : Loss Test = {loss_total_test:.4}')
              print()
```

Epoch = 0 : Loss Train = 3.829

Epoch = 0 : Loss Test = 2.482

Epoch = 10 : Loss Train = 0.6239
Epoch = 10 : Loss Test = 0.6271

Epoch = 20 : Loss Train = 0.5574
Epoch = 20 : Loss Test = 0.5661

Epoch = 30 : Loss Train = 0.5113
Epoch = 30 : Loss Test = 0.5209

Epoch = 40 : Loss Train = 0.4796
Epoch = 40 : Loss Test = 0.4892

Epoch = 50 : Loss Train = 0.4586
Epoch = 50 : Loss Test = 0.4686

Epoch = 60 : Loss Train = 0.4437
Epoch = 60 : Loss Test = 0.4538

Epoch = 70 : Loss Train = 0.4321
Epoch = 70 : Loss Test = 0.4421

Epoch = 80 : Loss Train = 0.4226
Epoch = 80 : Loss Test = 0.4324

Epoch = 90 : Loss Train = 0.4144
Epoch = 90 : Loss Test = 0.4243

Epoch = 100 : Loss Train = 0.4073
Epoch = 100 : Loss Test = 0.4174

Epoch = 110 : Loss Train = 0.4009
Epoch = 110 : Loss Test = 0.4108

Epoch = 120 : Loss Train = 0.395
Epoch = 120 : Loss Test = 0.4051

Epoch = 130 : Loss Train = 0.3896
Epoch = 130 : Loss Test = 0.3996

Epoch = 140 : Loss Train = 0.3847
Epoch = 140 : Loss Test = 0.3944

Epoch = 150 : Loss Train = 0.3801
Epoch = 150 : Loss Test = 0.3899

Epoch = 160 : Loss Train = 0.3759
Epoch = 160 : Loss Test = 0.3857

Epoch = 170 : Loss Train = 0.3719
Epoch = 170 : Loss Test = 0.3818

Epoch = 180 : Loss Train = 0.3685
Epoch = 180 : Loss Test = 0.3784

Epoch = 190 : Loss Train = 0.3651
Epoch = 190 : Loss Test = 0.3754

Epoch = 200 : Loss Train = 0.3622
Epoch = 200 : Loss Test = 0.3728

Epoch = 210 : Loss Train = 0.3595
Epoch = 210 : Loss Test = 0.3704

Epoch = 220 : Loss Train = 0.3569
Epoch = 220 : Loss Test = 0.368

Epoch = 230 : Loss Train = 0.3543
Epoch = 230 : Loss Test = 0.3657

Epoch = 240 : Loss Train = 0.352
Epoch = 240 : Loss Test = 0.3638

Epoch = 250 : Loss Train = 0.3499
Epoch = 250 : Loss Test = 0.3614

Epoch = 260 : Loss Train = 0.3477
Epoch = 260 : Loss Test = 0.3597

Epoch = 270 : Loss Train = 0.3457
Epoch = 270 : Loss Test = 0.3577

Epoch = 280 : Loss Train = 0.3438
Epoch = 280 : Loss Test = 0.356

Epoch = 290 : Loss Train = 0.342
Epoch = 290 : Loss Test = 0.3543

Epoch = 300 : Loss Train = 0.3401
Epoch = 300 : Loss Test = 0.3529

Epoch = 310 : Loss Train = 0.3386
Epoch = 310 : Loss Test = 0.3516

Epoch = 320 : Loss Train = 0.337
Epoch = 320 : Loss Test = 0.3501

Epoch = 330 : Loss Train = 0.3354
Epoch = 330 : Loss Test = 0.3489

Epoch = 340 : Loss Train = 0.3339
Epoch = 340 : Loss Test = 0.3475

Epoch = 350 : Loss Train = 0.3325
Epoch = 350 : Loss Test = 0.3462

Epoch = 360 : Loss Train = 0.3311
Epoch = 360 : Loss Test = 0.345

Epoch = 370 : Loss Train = 0.3298
Epoch = 370 : Loss Test = 0.3438

Epoch = 380 : Loss Train = 0.3284
Epoch = 380 : Loss Test = 0.3429

Epoch = 390 : Loss Train = 0.3271
Epoch = 390 : Loss Test = 0.3413

1.10 Evaluation

```
[20]: yp_total = []  
      yt_total = []  
      with torch.no_grad():  
          for x, y in test_loader:  
              yp = model(x)  
              yp_total.append(yp.squeeze())  
              yt_total.append(y)
```

```
[21]: yp_total = torch.cat(yp_total)  
      yt_total = torch.cat(yt_total)
```

```
[22]: r2_score(yp_total, yt_total)
```

```
[22]: 0.6572521143147981
```